ECE 777AB Fall 2023	Report in IEEE Format Group#: 3	ECE/WSU DRZ

Storage Media Trend Failure Analysis based on Enterprise workloads and environments

Note: Individual grades will be determined based on individual contributions!

Student Name (First Last)	Major Contributions
Sean Cowley; V688J944	33% - Data Extrapolation; Application, Results, Conclusions, Background Research
Yoel Woldeyes; J479Q742	33% - Theoretical Background, Methodology, Trials, Motivation
Md Reza E Rabbi; R844H794	33% - Abstract, Introduction, Literature Review
Optional Comments:	

***Abstract*—This paper addresses the growing complexities and risks in data storage systems in our data-centric world, focusing on the limitations of conventional methods for predicting storage media failure. These traditional methods, based on statistical analysis of past failures, often neglect various factors impacting storage media lifespan, workload characteristics, and environmental conditions. We propose an approach using Machine Learning, specifically Long Short-Term Memory (LSTM) networks within the TensorFlow package, and Decision Forests within the SciKit-Learn package, utilizing Self-Monitoring, Analysis, and Reporting Technology (SMART) data. This approach offers a promising solution for early detection of storage media failure. Parameters such as temperature, read/write errors, and power consumption, detectable using SMART data, are observed. The paper reviews the potential of LSTM and Decision Forests in analyzing SMART data to identify potential failures and failing indicators, enabling users to proactively maintain their systems. Our analysis leverages data center data packages publicly provided by Backblaze, a Data Storage company.**

I. Introduction

In today's data-driven world, the reliability, dependency, and efficiency of data storage systems are skyrocketing. Modern businesses heavily rely on data to execute essential applications, promote innovation, and provide a competitive advantage. However, due to the exponentially growth of data and increasing the number of variety and complexity, there is a significant danger and risk to corrupt data which increases the resource cost of operation. Conventional methods for storage media failure have been depending on statistical analysis of past failures. There are several limitations like unexpected downtime and data loss to this approach to find out the "Annualized Failure Rate" (AFR). It is based on historical data of Failed Storage Media in conjunction to the total quantity of operational storage media, it is very common that it ignores the diverse factors that influences storage media lifespan such as workload characteristics, environmental conditions, and operating mechanics. Machine Learning and Self-Monitoring, Analysis, and Reporting Technology (SMART) data offer a promising method for early storage media failure. Parameters like temperature, read/write errors, power consumption are few that may be found by using SMART data showing the effectiveness and condition of storage devices. Utilizing Machine Learning can help analyze SMART data to identify potential failures and indicators as well as promote data holders to take proactive action before difficulties occur. To analyze those issues, we are using some data center failure reports provided by Backblaze.

II. Literature Review

Hard Disk Drive (HDD) is a type of rotational storage media. It is a flat, round plate with a coating of magnetic substance made of glass and aluminum, known as a platter. The platter is rapidly rotating, with a head slightly above on a very narrow track. It will be difficult to read or write the data if there is even the slightest angular position error. Shock and Vibration to this arm can cause the head's accuracy to degrade resulting in degraded performance and missed sectors initializing failure states.

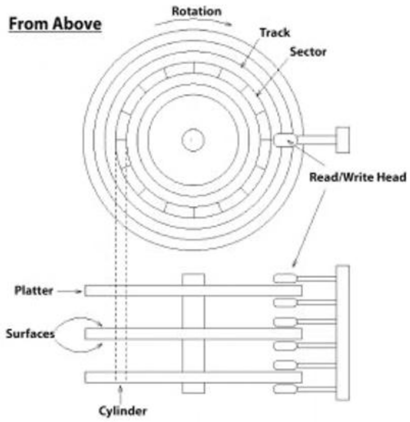

Figure 1: HDD Structure

The delay time will then rise as it tries to read/write the data again. Vibrations can occur for a variety of reasons, though a largely misrecognized source is the enclosure of the storage media. A series of HDD within the same enclosure can harmonize a detrimental frequency degrading performance.

Z. He; V. Venkataramanan; C.Y.Ng; E. H. Ong, used the LDV (Laser Dopler Vibrometer) and ALO (All in One) testers, and the results show that air pressure significantly affected how HDD components vibrated. When the air pressure dropped by 400 mbar during the read/write test, the hard drive exhibited performance issues and scratched media. but performed properly at 466 mbar. During that experiment, they also observed that, when air pressure dropped to 600 mbar, NRRO dramatically increased, which led to a fall in air density and viscosity. These increased the vibration's amplitude. [2]

A fan is required for cooling; however it can reduce HDD performance. Takehiko Eguchi and Kazuhide Ichikawa used the L8 array test with four factors to figure out the vibration. They investigate two types of vibration: structure-borne and air-borne. They discovered that the airborne conducted 80% of the power of PES. They proposed a chassis design and damping ratio to suppress that vibration [3]. In one study found that data center environment has a huge impact on the HDD's performance. In the data center environment, often time, temperature rise for several reason. The author came up with a solution with a design of data center which was able to reduce the temperature of that environment and increase the performance of HDD. [4] In another study, one author proposed a new reliability model for hard disk drive that considers head-disk separation as a critical factor. The model encompasses a range of environmental and operational parameters that affect HDD reliability. Here, the author highlighted the importance of maintaining an appropriate clearance between the magnetic recording head and the disk for mechanical reliability. [5]. the problem of long-tail delay in Solid-State Drives (SSDs) based on NAND flash by proposing a novel method that maximizes garbage collection by utilizing both short-term and long-term idle time periods. The authors point out that there is a chance to maximize garbage collection during idle time, and they suggest using idle time detectors to find these times. Their study offered a promising solution to improve SSD performance and suggested potential application for further enhancements. One study examines the correlation between temperature and HDD failures by analyzing data from a large population of systems within a data center. The study finds that temperature is correlated with HDD failures at three levels: within an enclosure, across a rack, and across multiple racks. The authors use an Arrhenius model to develop a reliability model that can be used to estimate AFR and select efficient data center operating temperatures. [7-8]. In another study, key SSD & NAND component reliability issues are examined to prove effects of drive cycling with the associated thermal shock, NAND component read disturbs, NAND "XOR" (internal RAID), and SSD mechanical design reliability issues [9]. The study finds that temperature is correlated with HDD failures at three levels: within an enclosure, across a rack, and across multiple racks [10]. Self-Monitoring, Analysis, and Reporting Technology (S.M.A.R.T.) is used extensively in this study with the Media Wearout Indicator used as a prime indicator of failure. End-Of-Life NANDs are often used to simulate a drive that has been used heavily within an intensive number of "Drive Writes Per Day". This figure is usually represented over a five-year time frame, once this number of DWPD multiplied by 265 by five, you can calculate the expected capacity overwrites a device can endure within a manufacturer's design specifications. [11-12]. Temperature has a severe effect on HDD performance. In one study, found out that AFR can double when temperature raised from 40 deg. Celsius to 55 deg. Celsius. Here, they use Baseboard Management Controller (BMC) to monitor the temperature across all components. The BMC controls fan speeds, adjusting the airflow within a server's chassis to provide optimal operational temperatures.

[13]. As we discussed before that vibration can disrupt HDD's performance. One author used Servo Mechanism which was used for feedback control loop position for reading/writing. This design compensates mechanical interference that can arise from vibration of fans and other systems in HDD. [14]. HDD failure rates vary across the industry and do not always follow the traditional bathtub curve. One study found that AFR for HDDs ranged from 1.7% in the first year of operation to 8.6% in year three and 7% in year five. Another study observed that HDDs exhibit an increasing-decreasing-stabilized pattern of failure rates during early deployments. These findings suggest that HDD failure rate modeling and field behaviors are diverse and require careful consideration. [15-16]. XGBoost metrics can be used to monitoring the health of HDD's performance. In one study, they compare healthy and unhealthy HDD. They observed that the median workload for unhealthy HDDs is 1.5x greater than the median workload for healthy HDDs. Age has also impact on HDD performance. In this paper they also analyze, normalize AFR is increased from 0 months to 10 months. And stable until 24 months. And after that AFR increased. But this may not be applicable for other models. The authors' environment is different from the typical bathtub curve for hardware failure rates. They don't see the high initial failure rate because most early life failures have already been screened out by the manufacturers. Instead, their failure curves start with an increase and then reach a steady state. This suggests that

there may be other factors at play in their environment that are causing the failure rates to increase over time [17]. In one study, author proposes a proactive solution to managing large-scale production storage systems. The goal is to predict disk failures before they occur and prevent data loss. The approach relies on SMART data and statistical analysis techniques to identify the start of disk degradation. The authors develop models to quantify disk health degradation and propose remediation mechanisms to prolong disk lifetime [18]. Now-a-days, Notebook computers, particularly their hard drives, are prone to impact-induced failures and vibration problems. Shock and vibration performance is critical during the product design phase to solve these concerns. To reduce the danger of failure, the author suggested, cushioning materials are routinely used. The dynamic properties and shock response of a system equipped with nonlinear rubber mounts are analyzed in this research using modal parameters derived from vibration testing. A better design for the rubber mounts is also proposed. [19]. External disturbance can hamper HDD’s performance. The author proposed a novel control method for HDD that utilizes an add-on nonlinear feedback controller to enhance disturbance rejection capabilities. The suggested technique improves on the existing linear controller by handling external disturbances well, especially when tracking errors become large. The suggested scheme's stability has been rigorously demonstrated, and experimental results from chassis tests on various market notebook PCs demonstrate the method's usefulness in handling wide-band vibration. Notably, the proposed scheme's implementation is straightforward and maybe easier than the linear controller [20]. Spindle motor performance has a direct impact on important performance and reliability issues like as rotational latency, drive access time, and spin-up time. The spin-up time is a crucial factor in power management strategies since it dictates how soon the disk platters accelerate from a motionless start to their maximum operating speed. Faster spin-up times are demanded by users, however this results in higher supply voltage, beginning current, and running current, which raises power consumption. Concern over power usage is growing, particularly for laptops and PCs with several storage devices. The author addressed those challenge of optimizing spin-up time while adhering to constraints such as variable rotational inertia, reduced power usage, and limited supply voltage and beginning current. For this performance optimization job, a computational model that combines a closed-loop feedback control technique with time-stepping FEM has been constructed [21].

III. BACKGROUND RESEARCH

Hard Disk Drives (HDDs) primarily fail due to their mechanical nature and the wear and tear that comes with it. The HDD is composed of spinning disks, or platters, and a read/write head that moves across these platters to access data. Over time, these mechanical components can wear out or break, leading to read/write errors or a complete inability to access the stored data. This mechanical failure is the most common reason for HDD failure.

Excessive heat is another significant factor that can lead to HDD failure. The heat can damage the electronic components of the drive or cause the mechanical parts to expand and malfunction. Therefore, it's crucial to keep HDDs in a cool, well-ventilated environment to prevent overheating.
Power surges or failures represent another common cause of HDD failure. Sudden changes in power can damage the electronic components of the HDD, leading to data corruption or drive failure. This is why it's often recommended to use surge protectors and uninterruptible power supplies (UPS) with computers that have HDDs.

Physical shock or damage is also a common cause of HDD failure. HDDs are sensitive to physical shock or vibration, and a sudden jolt or drop can cause the read/write head to touch the platters, leading to a "head crash." This can damage the platters, causing data loss and potentially rendering the drive unusable.

Manufacturing defects can also lead to HDD failure. These defects might not be apparent immediately but can cause the drive to fail prematurely. This is why even new drives can sometimes fail unexpectedly.

Finally, the corruption of the HDD's firmware, which controls the operation of the drive, can cause the drive to fail. The firmware can become corrupted due to software bugs, power surges, or malicious software, causing the drive to behave erratically or fail completely.

Solid State Drives (SSDs) fail for different reasons than Hard Disk Drives (HDDs) due to their distinct technology. Unlike HDDs, SSDs have no moving parts, so mechanical failure is not a concern. However, they have their own unique set of failure modes.

One of the primary reasons for SSD failure is wear leveling. SSDs store data in flash memory cells, and each cell can only endure a limited number of program/erase (P/E) cycles before it becomes unreliable. Wear leveling is a technique used to extend the lifespan of the SSD by distributing writes evenly across all the cells. However, over time, as the P/E cycles are exhausted, the cells start to fail, leading to data loss and eventually drive failure.

Another common cause of SSD failure is the failure of the power circuit. SSDs require precise control of power to program the flash memory cells. If there's a power surge or a failure in the power circuit of the SSD, it can lead to over-voltage or under-voltage conditions. These can damage the flash memory cells, leading to data corruption or complete drive failure.

Firmware corruption is another reason for SSD failure. The firmware of an SSD controls its operation, including the wear leveling algorithm and the translation between logical block addresses and physical flash memory cells. If the firmware becomes corrupted due to a bug, a power surge, or malicious software, it can cause the SSD to behave erratically or fail completely.

Manufacturing defects can also cause SSDs to fail. These defects can be in the flash memory cells themselves, the power circuit, or other components of the SSD. As with HDDs, these defects may not be apparent immediately but can cause the drive to fail prematurely.

Data corruption can also lead to SSD failure. This can occur due to software bugs, power surges, or hardware errors. Once the data is corrupted, it may not be possible to read it, and in severe cases, the SSD may not be able to function at all.

IV. MOTIVATION

The advent of Artificial Intelligence (AI) and its applications across various sectors has led to an unprecedented surge in data generation. This data is vital for training, validating, and testing complex AI models that drive innovation and efficiency in numerous fields. Consequently, there is an exponential demand for robust and reliable data storage solutions, particularly in large-scale data centers. These data centers serve as the backbone of the digital economy, storing vast amounts of information with a multitude of storage media technologies provided by various vendors.

However, with increasing reliance on these data repositories comes the challenge of maintaining storage media health to ensure data integrity and accessibility. The failure of storage media not only leads to data loss but can also cause significant downtime, leading to financial and reputational damage for businesses. Therefore, it is imperative to have predictive mechanisms in place that can forecast potential storage media failures, allowing for preemptive measures to be taken.

Current legacy methods of predicting storage media lifespan, such as Annualized Failure Rate (AFR), provide a reactive approach to failure. They often fail to capture the nuanced and complex patterns that precede a storage media's decline. Recent advances in Machine Learning (ML) have opened up new frontiers in predictive analytics, enabling more sophisticated and accurate prediction models. In particular, the analysis of Self-Monitoring, Analysis, and Reporting Technology (SMART) data through ML algorithms presents an opportunity to revolutionize how data centers manage and maintain storage media health.

The motivation for this research stems from the need to address these challenges head-on by developing ML-based predictive models that can detect early signs of storage media failures. By processing and analyzing SMART data from large-scale data centers, such as the comprehensive datasets provided by Backblaze, this research aims to uncover trends and indicators of failing storage media. The anticipated outcome is a model that can not only predict failures but also provide insights into the root causes, enabling more informed maintenance strategies.

This research is also motivated by the need to enhance the efficiency of data centers, reducing the frequency of reconstructions, replacements, and associated downtime. With the proposed methodology, we expect to deliver a model that will serve as a crucial tool for data center administrators, improving operational efficiency and reliability. Moreover, the insights gained from this study will contribute to the design of more resilient storage systems, capable of handling the increasing demands of the AI-driven digital landscape. The motivation behind this research is to leverage ML to deliver predictive insights into storage media health, thereby mitigating the risks of storage media failures in data centers. This approach is not only timely but essential in supporting the growth and sustainability of AI applications and the data centers that underpin them.

V. THEORETICAL BACKGROUND

In developing a predictive model for hard drive failures using SMART data, several classification techniques can be employed. These techniques are instrumental in categorizing the health status of hard drives into classes such as 'Optimal', 'Warning', and 'Critical'. Here, we delve into the theoretical aspects of the classification techniques that could be utilized in this research.

- Logistic Regression: A statistical method for analyzing a dataset in which there are one or more independent variables that determine an outcome. The outcome is measured with a dichotomous variable (in which there are only two possible outcomes). It is used for binary classification and provides probabilities that a given instance falls into one of the two categories.

- Decision Trees: A decision tree is a flowchart-like structure in which each internal node represents a "test" on an attribute, each branch represents the outcome of the test, and each leaf node represents a class label. The paths from root to leaf represent classification rules.

- Random Forests: An ensemble learning method for classification that operates by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes (classification) of the individual trees. Random forests correct for decision trees' habit of overfitting to their training set.

- Support Vector Machines (SVM): A set of supervised learning methods used for classification, regression, and outliers’ detection. An SVM model is a representation of the examples as points in space, mapped so that the examples of the separate categories are divided by a clear gap that is as wide as possible.

- Neural Networks: A neural network is a network or circuit of neurons, or in a modern sense, an artificial neural network, composed of artificial neurons or nodes. For classification, a neural network will learn the input patterns that correspond to each class through training and then predict the class of new inputs based on this learned pattern.

- Long Short-Term Memory Networks (LSTMs): A special kind of recurrent neural network capable of learning long-term dependencies. LSTMs are well-suited to classifying, processing, and making predictions based on time series data, hence, they are a fitting choice for modeling time-dependent patterns in SMART data for predicting hard drive failures.

- K-Nearest Neighbors (KNN): A non-parametric method used for classification and regression. A sample is classified by a plurality vote of its neighbors, with the sample being assigned to the class most common among its k nearest neighbors (k is a positive integer, typically small).

- Naive Bayes Classifiers: A family of simple "probabilistic classifiers" based on applying Bayes' theorem with strong (naïve) independence assumptions between the features. They are among the simplest Bayesian network models but coupled with kernel density estimation, they can achieve higher accuracy levels.

- Gradient Boosting Machines (GBM): A machine learning technique for regression and classification problems, which produces a prediction model in the form of an ensemble of weak prediction models, typically decision trees.

Each of these techniques has its strengths and weaknesses, and often the best approach is to combine multiple models in an ensemble method to improve the predictive performance and robustness of the solution. This research will focus on exploring these techniques and identifying the most effective approach for predicting hard drive failures. We will be using the LSTMs and decision forests for this project.

VI. Methodology

This project outlines a comprehensive process for analyzing Self-Monitoring, Analysis, and Reporting Technology (SMART) data with the objective of predicting hard drive failures. The process is designed to leverage machine learning techniques and advanced analytics to detect trends in failing models and provide root cause analysis. The methodology is divided into eight key stages: data collection, preprocessing, feature extraction, SMART weight application, TensorFlow model development, prediction model training, testing, and outcome interpretation.

The data for this process will be sourced from Backblaze's published Drive Stats, and the TensorFlow framework with Python will be utilized for model applications. The goal is to identify correlated trends of failing indicators across logged features, and to pinpoint the worst-case drive model based on the data. Each stage of the process is designed to refine the data and enhance the predictive capabilities of the model, ensuring that the final output provides actionable insights for maintaining the health of storage devices in a data center.

1. Training SMART Data: This initial phase will involve the collection of SMART data from a diverse set of storage devices. The collected data, which includes various attributes that potentially indicate the health of a drive, will be curated to ensure a comprehensive dataset representing different manufacturers and usage conditions. Special attention will be given to the temporal aspects of data to account for the aging factor of the devices.
2. Preprocess: The preprocessing step will consist of several data cleaning sub-steps, such as outlier removal and data imputation, to handle anomalies and missing values. We will also standardize and normalize the metrics to a uniform scale to prevent any single metric from disproportionately influencing the prediction outcome due to variance in measurement scales.
3. Feature Extraction: In this stage, advanced analytical techniques, potentially including machine learning algorithms like decision trees, will be employed to identify and extract the most critical SMART attributes that correlate with drive failures. We will also explore feature engineering to create new predictive variables from the existing data.

4. SMART Weight Application: We will apply weights to the extracted features based on their predictive importance, determined through techniques such as correlation analysis and feature importance ranking from ensemble learning models. This will optimize the data input for the predictive model to focus on the most significant predictors.

5. TensorFlow Model Development: A specific TensorFlow model architecture will be selected or designed to accommodate the nature of the SMART data. Replacing the TBA TensorFlow model, we will design and implement a Long Short-Term Memory (LSTM) network. LSTMs are a type of recurrent neural network (RNN) suitable for making predictions based on time-series data, like SMART metrics, which are inherently sequential and temporal in nature. The LSTM model is capable of learning long-term dependencies and patterns in the data which are essential for anticipating hard drive failures.

6. Prediction Model: This model will be trained using the weighted features to forecast potential drive failures, classifying the health status of the drives into 'Optimal', 'Warning', and 'Critical'. Model selection will be based on a balance of accuracy and computational efficiency, considering real-world application constraints.

7. Test SMART Data: The robustness and accuracy of the prediction model will be tested using a separate dataset of SMART data that was not exposed to the model during the training phase. This will allow us to measure the model's performance and generalize its predictive capability on unseen data, which is critical for assessing its real-world applicability.

8. Outcome Interpretation (Optimal, Warning, Critical): The model's output will be interpreted to facilitate actionable insights. Drives marked 'Optimal' will be considered in good health, while 'Warning' will indicate the need for

closer monitoring, and 'Critical' will signal an immediate risk of failure, necessitating urgent backup and potential drive replacement.

The flowchart below represents a process for analyzing SMART (Self-Monitoring, Analysis, and Reporting Technology) data with the aim of predicting hard drive failures.

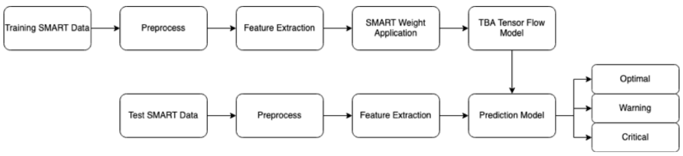

Figure 2: Process flowchart

The goal is to process failure analysis data from a data center's storage media to detect trends in failing models and provide a root cause analysis. As we mentioned above the source of training, validating, and testing data will be Backblaze's published Drive Stats, and we will use TensorFlow with Python for model applications. The expected outcomes are to find a correlated trend of failing indicators across logged features and to identify a worst-case drive model with indicating data.

VII. Trials

Legacy methods of determining a vintage of storage media had prior been a reactive action to a series of statistical analysis. A primary method had been Annualized Failure Rate, to rate a specific model of drive to a failure metric based on quantity and time on hours. New trends of Machine Learning have increased the potential of detecting these undesirable traits through analysis of Self-Monitoring, Analysis, and Reporting Technology SMART Logging. This allows for an increase of efficiency due to fewer reconstructions, replacements, and downtime.

In this section, we divide the content into four parts. The first part is preparing the data, the second part is doing the analysis, the third part is training the model, and the fourth part is evaluating and comparing the results.

1. Data Collection and Preparation: We gathered SMART data from Backblaze's published Drive Stats, focusing on a set of features such as date, serial number, manufacturer_model, capacity, failure boolean, and location, alongside SMART attributes 0-255. The dataset consisted of a total of 142,838 drives, out of which 6 had failed. This information was crucial as it provided a balanced perspective of both functioning and failed drives. The data underwent preprocessing to ensure consistency and usability for model training and testing.
2. Exploratory Data Analysis (EDA): An initial EDA will be conducted to understand the distributions, correlations, and patterns within the data. This step is crucial for feature selection and to determine the preprocessing requirements such as normalization and handling of missing data.
3. Model Training: Various machine learning models will be trained using preprocessed data. This will include LSTM and Decision forests. The LSTM model will be particularly focused on due to its ability to process time-series data, which is a significant component of SMART attributes.
4. Model Evaluation: The models' performances will be evaluated using metrics such as accuracy, precision, recall, and F1 score. For the LSTM model, we will also consider its ability to capture long-term dependencies in the SMART data for predicting failures.

VIII. Data Extrapolation

SMART, which stands for Self-Monitoring, Analysis, and Reporting Technology, is a monitoring system included in hard disk drives (HDDs), solid-state drives (SSDs), and eMMC drives. Its primary function is to detect and report various indicators of drive reliability with the intent of anticipating hardware failures. SMART values are a set of specific, pre-defined attributes or parameters that a drive tracks and reports. These values can provide information about various aspects of a drive's performance and health. These SMART values have intended defined attributes but are interpreted by each manufactuer differently.

Examples of SMART values and their intended descriptions include:

Read Error Rate: This indicates the rate of hardware read errors that occurred when reading data from a disk surface.

Reallocated Sectors Count: This represents the count of reallocated sectors. When the hard drive finds a read/write/verification error, it marks this sector as "reallocated" and transfers data to a special reserved area (the spare area).

Power-On Hours (POH): This is a count of the total power-on time of the hard drive, often reported in hours.

Temperature: This is the current temperature of the hard drive.

Write Error Rate: This is similar to the Read Error Rate but pertains to write operations.

Spin-Up Time: This measures the time it takes for the drive to spin its disk up to operational speed from a standstill condition. If this value is high, it could indicate problems with the drive spinning up.

Start/Stop Count: This value indicates the number of start/stop cycles of the hard drive. Drives with a high count might be older or heavily used.

Seek Error Rate: This value represents the rate of errors occurred while positioning the hard drive's heads (i.e., when the drive is seeking a specific position on the disk). A high rate could indicate a mechanical problem.

Spin Retry Count: This value indicates the count of retry of spin start attempts. This attribute stores a total count of the spin start attempts to reach the fully operational speed (under the condition that the first attempt was unsuccessful). An increase of this attribute value is a sign of problems in the hard disk mechanical subsystem.

Airflow Temperature: This value indicates the air temperature flowing past the drive. High temperatures could indicate inadequate cooling.

Command Timeout: This value indicates the number of aborted operations due to HDD timeout. Usually, this attribute is a sign of problems in the hard disk power supply circuitry.

Reallocation Event Count: This value indicates the total count of attempts to transfer data from reallocated sectors to a spare area. Both successful and unsuccessful attempts are counted.

The results of these SMART values are used to predict the likelihood of a drive failure. For example, a high number of reallocated sectors could indicate a failing drive. By monitoring these values, system administrators can identify problematic drives and replace them before they fail, preventing data loss and maintaining system performance.

It's important to note that while SMART can help anticipate some disk failures, it is not always accurate and some failures can occur without any predictive SMART values. Therefore, it should be used as a part of a comprehensive data protection strategy, not as the sole method of anticipating drive failures. To expand upon this paper, an effective way of applying these SMART values would be to weigh each attribute depending on the risk introduced to the storage media by abnormal values.

IX. Application

Building a failure prediction model or tree using SMART (Self-Monitoring, Analysis, and Reporting Technology) values involves a series of steps. Initially, data collection is crucial, which involves gathering SMART values from a large number of drives over a period of time. This dataset should ideally include both drives that have failed and those that have not, to provide a balanced perspective. Each drive is then labeled based on whether it failed within a certain time frame after each SMART value was recorded. This becomes the target variable that the model aims to predict. Not all SMART values are equally useful for predicting drive failure, so statistical techniques are employed to identify which SMART values are most strongly associated with drive failure, a process known as feature selection. This data is all available within the Backblaze dataset providing the SMART values correlated to a single drive's Serial Number once per day.

Once the relevant features are selected, a machine learning algorithm is used to train a model on the labeled data. This could be a classification model if you're predicting whether a drive will fail or not, or a regression model if you're predicting the time until failure. The model is then validated by splitting the data into a training set and a validation set, training the model on the training set, and testing its performance on the validation set. This ensures the model can generalize to new data. Once validated, the model can predict drive failures based on SMART values. If a drive's predicted probability of failure exceeds a certain threshold, it could be flagged for replacement. The model should be continuously updated with new data to keep it accurate and up-to-date. This process enables the prediction of drive failures before they happen, allowing for data to be backed up and drives to be replaced before a failure occurs.

To fulfill this objective, we employ a combination of Python, TensorFlow, and SciKit-Learn, which are powerful tools for data analysis and machine learning. Python serves as our primary programming language due to its simplicity, versatility, and the robustness of its scientific computing and data analysis libraries. TensorFlow, an open-source library developed by Google, is used for building machine learning models, specifically Long Short-Term Memory (LSTM) networks in this case. LSTM networks are a type of recurrent neural network that are particularly adept at processing sequential data, making them ideal for analyzing the time-series data represented by SMART values. On the other hand, we use SciKit-Learn, another open-source library that provides a wide range of tools for machine learning and predictive data analysis, to construct Decision Forests.

Decision Forests, also known as Random Forests, are an ensemble learning method that operates by constructing multiple decision trees and outputting the class that is the mode of the classes or mean prediction of the individual trees. By leveraging both Decision Forests and LSTM networks, we aim to create a comprehensive, robust model capable of accurately identifying failure indicators in hard drives. This combination allows us to capture both the intricate time-dependent patterns in the SMART values, as well as the more straightforward, static relationships between different SMART values and drive failure.

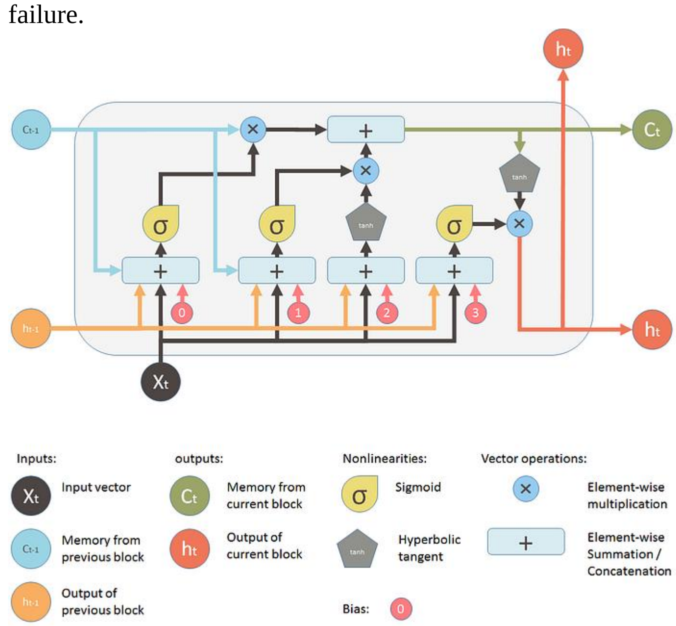

Figure 3: LSTM Model Flowchart

The decision forest application was accomplished with the following methodology. We begin by importing the necessary libraries for data manipulation, machine learning, and data visualization, such as pandas, TensorFlow, scikit-learn, and matplotlib.

To load the data, we use the glob library to find all CSV files in a specified directory. Each of these files is then read into a pandas DataFrame. We store these individual DataFrames in a list and concatenate them into a single DataFrame, df.

Next, we prepare the data by separating the target label 'failure' from the features in the DataFrame. We store the features in X and the target label in y. Following that, we remove the 'date', 'serial_number', and 'model' columns from X, as these are not likely to be useful for predicting drive failure.

Once the data is prepared, we split it into a training set and a test set using the train_test_split function from scikit-learn. We use 80% of the data for training and reserve 20% for testing.

With a training set and a test set ready, we create a Decision Tree Classifier and train it on the training data using the fit method. After the model is trained, we use it to make predictions on the test set.

To evaluate the performance of the model, we compute the accuracy of the predictions by comparing them to the actual labels of the test set. This gives us a quantifiable measure of how well our model is performing.

Finally, we visualize the decision tree using matplotlib and scikit-learn's plot_tree function. This plot shows the decisions made at each node in the tree, providing a visual representation of how the model makes its predictions.

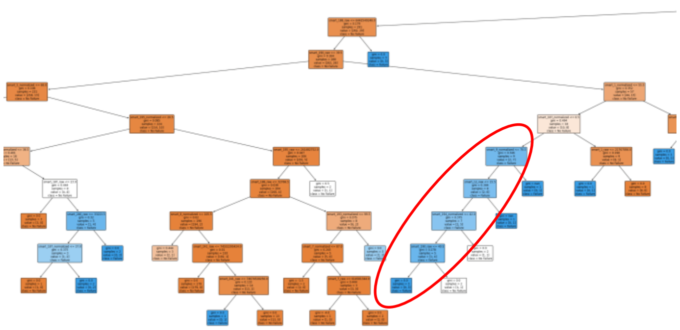

Figure 4: Decision Forest

In the Model application of our study, we utilized various Python libraries to facilitate the process. We used pandas for data manipulation, TensorFlow for building and training the LSTM model, scikit-learn for data preprocessing and model evaluation, and matplotlib for data visualization.

The first step in our application was to load the data. We utilized the glob library to identify all CSV files in a specified directory, each of which was read into a pandas DataFrame. These individual DataFrames were then concatenated into a single DataFrame for further processing.

Next, we performed a check for NaN and infinite values within the DataFrame. If any NaN values were identified, they were replaced with the mean of the respective column. Similarly, any infinite values were replaced with the maximum finite value of the respective column. This process was carried out only for numeric columns, as the mean and maximum finite value are not defined for non-numeric columns.

Following the data cleaning process, we selected our features and target from the DataFrame. The features included all columns except 'date', 'serial_number', 'model', 'capacity_bytes', and 'failure', while the target was the 'failure' column.

eeeeee

To ensure all features were on the same scale, we normalized them to fall within a range of 0 to 1 using the MinMaxScaler from scikit-learn. This step is crucial as many machine learning algorithms require input features to be on a similar scale for optimal performance.

We then split our data into a training set and a testing set using the train_test_split function from scikit-learn, with 80% of the data used for training and the remaining 20% for testing.

To prepare the data for the LSTM model, we reshaped the input to be 3D [samples, timesteps, features], as LSTM models in TensorFlow require input in this format.

We defined our LSTM model using the Sequential API from TensorFlow. The model consisted of an LSTM layer with 50 units and a 'relu' activation function, followed by a dense output layer with a single unit.

The user was then given the option to choose an optimizer for the model. Depending on the user's choice, the model was compiled with the Adam optimizer with gradient clipping, the RMSprop optimizer, or the Adam optimizer with a custom learning rate. The loss function was set to 'binary_crossentropy', suitable for binary classification problems, and accuracy was tracked as a metric during training.

The model was then trained on the training data for 200 epochs. After training, the model's performance was evaluated on the test set, providing us with the test loss and accuracy. The test loss, a measure of how well the model's predictions match the actual values, and the test accuracy, a measure of how often the model's predictions match the actual values, provided us with a quantifiable measure of our model's performance.

X. RESULTS

In our research, we applied a Decision Tree model over a calendar year quarter to analyze and predict hard drive failures based on Self-Monitoring, Analysis, and Reporting Technology (SMART) values. Our model identified a specific series of SMART values that consistently indicated an impending failure. These SMART values were 9, 12, 194, and 190.

SMART value 9, known as Power-On Hours (POH), indicates the total number of hours the hard drive has been powered on and is often used as a rough indicator of the age of the drive. SMART value 12, the Drive Power Cycle Count, indicates the count of full power-on to power-off cycles of the hard drive, with a high count potentially indicating frequent power cycles that may impact the drive's lifespan. SMART value 194, the Drive Temperature, indicates the current temperature of the hard drive in Celsius, where high or rapidly fluctuating temperatures could signal potential issues with the drive or its environment. Lastly, SMART value 190, the Airflow Temperature, represents the temperature of the air flowing past the hard drive in Celsius, with high or rapidly fluctuating values also serving as potential indicators of issues.

The emergence of these specific values as predictors of failure underscores the importance of monitoring multiple parameters to accurately assess drive health and predict failures. Our observations showed that a compound consideration of these values provided a significant indication of imminent drive failure, allowing for preventative actions to be taken. Please note that these are typical interpretations of these attributes, though the exact definitions can sometimes vary between different hard drive manufacturers.

In our application of the Long Short-Term Memory (LSTM) model for predicting hard drive failures, we observed some interesting results. The test loss, which is a measure of the model's error or how much the model's predictions deviate from the actual values, was reported as nan, an acronym for "Not a Number". This typically indicates an issue during the calculation of the loss, possibly due to numerical instability, issues with data preprocessing, or problems with the model's architecture or learning rate. This is a point of concern as the loss is a crucial metric for judging the model's performance, with lower values indicating better performance.

On the other hand, the test accuracy of the model was exceptionally high, approximately 99.996%. This high accuracy might be influenced by the imbalance in our dataset, which consisted of a total of 142,838 drives, out of which only 6 had failed. Given the significantly higher number of functioning drives compared to failed ones, the model might be correctly predicting the majority class (functioning drives) most of the time, leading to a high accuracy. However, it's important to note that in such imbalanced datasets, accuracy might not be the best metric to evaluate model performance, as it doesn't take into account how well the model is predicting the minority class (failed drives).

Given the nan loss and the extremely high accuracy, it is important for us to validate these results and ensure they are not due to factors like an imbalance in the target variable classes, overfitting to the training data, or potential errors in the data or model implementation. These results underscore the need for thorough model evaluation and validation in machine learning-based predictive maintenance.

XI. Limitations and Future Work

One of the main limitations of this study is the size of the dataset. As the dataset grows quarter over quarter, the total size is nearing 40GB. This large volume of data poses challenges in terms of computational resources and time required for model training. In our current setup, training the model on the entire dataset would take close to 100 continuous days, which is not feasible for real-time or near real-time predictive maintenance applications.

In future work, several strategies could be explored to address this limitation. One approach could be to use a subset of the data for model training, focusing on the most recent data or randomly sampling from the entire dataset. However, this could potentially miss important patterns in the data and reduce the predictive accuracy of the model.

Another approach could be to use more efficient models or optimization algorithms that can handle large datasets more effectively. This could involve exploring other machine learning algorithms that are known for their scalability and efficiency.

Additionally, the use of more powerful computational resources or distributed computing techniques could significantly reduce the training time. This could involve using high-performance computing clusters or cloud-based machine learning platforms that can distribute the computation across multiple machines.

Finally, techniques such as incremental learning or online learning could be explored. These techniques allow the model to learn from new data without having to retrain on the entire dataset, which could be particularly useful as new data becomes available each quarter.

By exploring these strategies, we aim to develop a more scalable and efficient predictive maintenance solution that can handle the growing volume of data in future work.

XII. Conclussion

In conclusion, our study, we found that specific SMART values, namely 9, 12, 194, and 190, consistently signaled an impending drive failure. SMART value 9, or Power-On Hours (POH), is a measure of the total number of hours the hard drive has been powered on, serving as a rough indicator of the drive's age. A high POH could indicate an older, potentially more worn-out drive, increasing the likelihood of failure.

SMART value 12, the Drive Power Cycle Count, measures the number of full power-on to power-off cycles of the hard drive. A high count could suggest frequent power cycles, which may stress the drive and shorten its lifespan.

SMART values 194 and 190, the Drive Temperature and Airflow Temperature respectively, indicate the current and airflow temperatures of the hard drive in Celsius. High or rapidly fluctuating temperatures could signal potential issues with the drive or its environment, such as inadequate cooling, which could lead to overheating and eventual drive failure.

These findings underscore the importance of monitoring multiple parameters to accurately assess drive health and predict failures. However, it's important to note that while these SMART values can help predict some disk failures, they are not always accurate and some failures might occur without any predictive SMART values. Therefore, SMART should be used as part of a comprehensive data protection strategy, not as the sole method of anticipating drive failures.

Our LSTM model achieved an exceptionally high test accuracy of approximately 99.996%, suggesting that the model is very effective at classifying whether a drive will fail or not based on the input SMART values. However, the test loss reported as 'nan' indicated potential issues that required further investigation. This could be due to numerical instability, issues with data preprocessing, or problems with the model's architecture or learning rate. This highlights the need for thorough model evaluation and validation in predictive maintenance.

Overall, our research showed the potential of SMART values in predicting drive failures and the effectiveness of machine learning models in analyzing these values. Our findings could contribute significantly to preventative measures in data protection and drive maintenance, improving system reliability and performance.

REFERENCES

[1] [https://www.ept.ca/features/everything-need-know-hard-drive-vibration/]

[2] Z. He, V. Venkataramanan, C. Y. Ng and E. H. Ong, "Vibration characteristics of hard disk drives with different internal air pressures," 2009 Asia-Pacific Magnetic Recording Conference, Singapore, 2009

[3] [T. Eguchi, K. Ichikawa, S. Takada and M. Takada, "High Frequency Dynamics of a HDD in a Storage Server: Airborne Vibration Transmission and PES Simulations," 2018 Asia-Pacific Magnetic Recording Conference (APMRC), Shanghai, China, 2018

[4] J. Zhao et al., "Disk Failure Early Warning Based on the Characteristics of Customized SMART", 2020 19th IEEE Intersociety Conference on Thermal and Thermomechanical Phenomena in Electronic Systems (ITherm), 2020

[5] K. -Y. Chen, C. -H. Wu and C. -T. Lee, "Short-Term and Long-Term Idle Time Detectors for Reducing Long-Tail Latency in Solid-State Drives," 2023 Sixth International Symposium on Computer, Consumer and Control (IS3C), Taichung, Taiwan, 2023

[6] B. D. Strom, S. C. Lee, G. W. Tyndall and A. Khurshudov, "Hard Disk Drive Reliability Modeling and Failure Prediction," *Asia-Pacific Magnetic Recording Conference 2006*, Singapore

[7]

[8] K. Aruga, M. Suwa, K. Shimizu, and T. Watanabe, "A study on positioning error caused by flow induced vibration using helium-filled disk drives", IEEE Trans. Magn., September 2007.

[9] T. Egnchi, "Aerodynamic damping on vibration of rotating disks and effect of shroud and top cover", IEEE Intermag 2008, Madrid, Spain.

[10] Mielke, N.; Marquart, T.; Ning Wu; Kessenich, J.; Belgal, H.; Schares, Eric; et al., "Bit error rate in NAND Flash memories," Reliability Physics Symposium, 2008. IRPS 2008. IEEE International, April 27 2008-May 1 2008

[11] Chen, F., Zou, D., & Sun, X. “Analyzing the Correlation between Temperature and Hard Disk Drive Failure in a Large-Scale Data Center" (2015)

[12] Yu Cai; Yalcin, G.; Mutlu, O.; Haratsch, E.F.; Cristal, A.; Unsal, O.S.; Ken Mai, "Flash correct-and-refresh: Retention-aware error management for increased flash memory lifetime," Computer Design (ICCD), 2012 IEEE 30th International Conference, Sept. 30 2012-Oct. 3 2012

[13] Xu, C., Chuang, Y., & Sun, Y. (2019). SMART-Aware Garbage Collection for Solid-State Drives in Real-Time Systems. IEEE Transactions on Computers, 68(10), 1454-1467.

[14] S. Sankar, M. Shaw and K. Vaid, "Impact of temperature on hard disk drive reliability in large datacenters", *2011 IEEE/IFIP 41st International Conference on Dependable Systems Networks (DSN)*, pp. 530-537, 2011.

[15] X. Chen, Hard Disk Drive Control Mechanical Systems Control Lab, Berkeley, 2016,

[16] Eduardo Pinheiro, Wolf-Dietrich Weber and Luiz André Barroso, "Failure trends in a large disk drive population", 2007.

[17] on G. Elerath and Sandeep Shah, "Server class disk drives: How reliable a re they?", *In Proceedings of the Annual Symposium on Reliability and Maintainability*, January 2004.

[18] Z. Miller, O. Medaiyese, M. Ravi, A. Beatty and F. Lin, "Hard Disk Drive Failure Analysis and Prediction: An Industry View," *2023 53rd Annual IEEE/IFIP International Conference on Dependable Systems and Networks - Supplemental Volume (DSN-S)*, Porto, Portugal, 2023

[19] S. Huang, S. Liang, S. Fu, W. Shi, D. Tiwari and H. -b. Chen, "Characterizing Disk Health Degradation and Proactively Protecting Against Disk Failures for Reliable Storage Systems," 2019 IEEE International Conference on Autonomic Computing (ICAC), Umea, Sweden, 2019

[20] S. Lim, Y. -B. Chang, N. -C. Park and Y. -P. Park, "Optimal Design of Rubber Mounts Supporting Notebook HDD for Shock and Vibration Isolation," *Asia-Pacific Magnetic Recording Conference 2006*, Singapore, 2006

[21] Q. Jia, S. Cao, G. Guo and J. Yu, "External Disturbance Rejection by Use of an Add-On Nonlinear Controller in HDD Servo Systems," in *IEEE Transactions on Magnetics*, vol. 49, no. 6, pp. 2624-2627, June 2013,

[22] Hla Nu Phyu, M. A. Jabbar and Z. J. Liu, "Computational optimization of the run-up performance of a HDD spindle motor," IEEE International Conference on Electric Machines and Drives, 2005., San Antonio, TX, USA, 2005

[23] https://www.backblaze.com/cloud-storage/resources/hard-drive-test-data

[24] https://blog.mlreview.com/understanding-lstm-and-its-diagrams-37e2f46f1714